

# revised-machine-learning-balaram-1

April 18, 2024

## 1 Week 5 Lab: Supervised Learning

This week's assignment will focus constructing and improving the performance of a KNN model.

### 1.1 Dataset:

**Dataset:** bank-additional-full.csv (Provided in folder assign\_wk5) ## Assignment Requirements

After reviewing the dataset's descriptive file (bank-additional-names.txt), define a research question to guide your analysis. Make sure your KNN Analysis addresses the following tasks: - Cleanup the dataset as you deem appropriate. As always, defend your reasoning!!! - Missing values? - Column names - Prepare the data for machine learning - A little EDA goes a long way - Do you need to do anything about data types? - KNN Analysis - What is your optimal K? - Evaluate the accuracy of your model - Discuss ways to improve the performance of your KNN model. \* Notice the requirement states **ways** - meaning more than one! \* Defend and backup your thoughts!!!!!! - KNN Model Improvement - Implement one of those methods to improve your KNN model performance. - Did your second model perform better than the first? - Conclusion/Summary - Compare and contrast your 2 models - Include numbers/graphs corresponding to your conclusions - Defend and backup your thoughts!!!!!!

## 2 Deliverables:

Upload your Jupyter Notebook to the corresponding location in WorldClass.

**Note::** Make sure you have clearly indicated each assignment requirement within your notebook.

**Important:** Make sure you have clearly indicated each assignment requirement within your notebook. Also, I highly encourage you to use markdown text to create a notebook that integrates your analysis within your code. The narrative within your notebook will count for 50% of your total grade on this assignment.

```
[1]: # Import essential data analysis libraries
import pandas as pd # For data manipulation and analysis
from sklearn.preprocessing import LabelEncoder # For converting text labels
    ↪ into a numeric form

# Visualization libraries to help us see trends and insights
import seaborn as sns # For making attractive and informative statistical
    ↪ graphics
```

```
import matplotlib.pyplot as plt # For creating static, interactive, and
↳ animated visualizations in Python

# Machine learning tools from scikit-learn
from sklearn.neighbors import KNeighborsClassifier # For using the K-Nearest
↳ Neighbors algorithm
from sklearn.model_selection import train_test_split # For splitting data into
↳ training and testing sets
from sklearn.pipeline import Pipeline # For sequential application of a list
↳ of transforms and a final estimator
from sklearn.preprocessing import StandardScaler # For scaling features to
↳ standardize the range of independent variables
```

**Problem statement :** Identify key factors that predict if a client will subscribe to a bank's term deposit program.

Load and view the data file

```
[2]: import pandas as pd

# Load the dataset
df = pd.read_csv('C:/Users/balar/Downloads/assign_wk5/assign_wk5/
↳ bank-additional-full.csv')

# Display the first few rows of the dataframe
df.head(20)
```

```
[2]: age;"job";"marital";"education";"default";"housing";"loan";"contact";"month";
"day_of_week";"duration";"campaign";"pdays";"previous";"poutcome";"emp.var.rate"
;"cons.price.idx";"cons.conf.idx";"euribor3m";"nr.employed";"y"
0 56;"housemaid";"married";"basic.4y";"no";"no";...
1 57;"services";"married";"high.school";"unknown"...
2 37;"services";"married";"high.school";"no";"ye...
3 40;"admin."; "married";"basic.6y";"no";"no";"no...
4 56;"services";"married";"high.school";"no";"no...
5 45;"services";"married";"basic.9y";"unknown";"..."
6 59;"admin."; "married";"professional.course";"n...
7 41;"blue-collar";"married";"unknown";"unknown"...
8 24;"technician";"single";"professional.course"...
9 25;"services";"single";"high.school";"no";"yes..."
10 41;"blue-collar";"married";"unknown";"unknown"...
11 25;"services";"single";"high.school";"no";"yes..."
12 29;"blue-collar";"single";"high.school";"no";"..."
13 57;"housemaid";"divorced";"basic.4y";"no";"yes..."
14 35;"blue-collar";"married";"basic.6y";"no";"ye..."
15 54;"retired";"married";"basic.9y";"unknown";"y..."
16 35;"blue-collar";"married";"basic.6y";"no";"ye..."
```

```

17 46;"blue-collar";"married";"basic.6y";"unknown...
18 50;"blue-collar";"married";"basic.9y";"no";"ye...
19 39;"management";"single";"basic.9y";"unknown";...

```

Use separator and view data frame

```

[3]: df = pd.read_csv('C:/Users/balar/Downloads/assign_wk5/assign_wk5/
↳bank-additional-full.csv', sep=';')
df

```

```

[3]:      age      job  marital      education  default  housing  loan  \
0      56  housemaid  married      basic.4y      no      no      no
1      57  services  married  high.school  unknown      no      no
2      37  services  married  high.school      no      yes      no
3      40   admin.  married      basic.6y      no      no      no
4      56  services  married  high.school      no      no      yes
...    ...
41183   73   retired  married  professional.course      no      yes      no
41184   46 blue-collar  married  professional.course      no      no      no
41185   56   retired  married  university.degree      no      yes      no
41186   44  technician  married  professional.course      no      no      no
41187   74   retired  married  professional.course      no      yes      no

```

```

      contact month  day_of_week  ...  campaign  pdays  previous  \
0  telephone   may      mon      ...      1      999      0
1  telephone   may      mon      ...      1      999      0
2  telephone   may      mon      ...      1      999      0
3  telephone   may      mon      ...      1      999      0
4  telephone   may      mon      ...      1      999      0
...    ...
41183  cellular  nov      fri      ...      1      999      0
41184  cellular  nov      fri      ...      1      999      0
41185  cellular  nov      fri      ...      2      999      0
41186  cellular  nov      fri      ...      1      999      0
41187  cellular  nov      fri      ...      3      999      1

```

```

      poutcome  emp.var.rate  cons.price.idx  cons.conf.idx  euribor3m  \
0  nonexistent      1.1      93.994      -36.4      4.857
1  nonexistent      1.1      93.994      -36.4      4.857
2  nonexistent      1.1      93.994      -36.4      4.857
3  nonexistent      1.1      93.994      -36.4      4.857
4  nonexistent      1.1      93.994      -36.4      4.857
...    ...
41183  nonexistent      -1.1      94.767      -50.8      1.028
41184  nonexistent      -1.1      94.767      -50.8      1.028
41185  nonexistent      -1.1      94.767      -50.8      1.028
41186  nonexistent      -1.1      94.767      -50.8      1.028

```

41187	failure	-1.1	94.767	-50.8	1.028
-------	---------	------	--------	-------	-------

	nr.employed	y
0	5191.0	no
1	5191.0	no
2	5191.0	no
3	5191.0	no
4	5191.0	no
...	...	...
41183	4963.6	yes
41184	4963.6	no
41185	4963.6	no
41186	4963.6	yes
41187	4963.6	no

[41188 rows x 21 columns]

Check shape, info, describe for our data frame

```
[4]: df.shape
```

```
[4]: (41188, 21)
```

```
[5]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 41188 entries, 0 to 41187
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   age                   41188 non-null  int64
1   job                   41188 non-null  object
2   marital               41188 non-null  object
3   education             41188 non-null  object
4   default               41188 non-null  object
5   housing               41188 non-null  object
6   loan                  41188 non-null  object
7   contact               41188 non-null  object
8   month                 41188 non-null  object
9   day_of_week           41188 non-null  object
10  duration              41188 non-null  int64
11  campaign              41188 non-null  int64
12  pdays                 41188 non-null  int64
13  previous              41188 non-null  int64
14  poutcome              41188 non-null  object
15  emp.var.rate          41188 non-null  float64
16  cons.price.idx        41188 non-null  float64
17  cons.conf.idx         41188 non-null  float64
```

```

18 euribor3m      41188 non-null float64
19 nr.employed    41188 non-null float64
20 y              41188 non-null object
dtypes: float64(5), int64(5), object(11)
memory usage: 6.6+ MB

```

The dataset doesn't have any blank spots, as we've seen from checking all the rows and columns with the `info()` method. However, the description tells us that missing information is actually labeled as 'unknown'. We should take a closer look at the data before we start cleaning it up.

```
[6]: df.describe()
```

```

[6]:
count    age      duration      campaign      pdays      previous \
count  41188.000000  41188.000000  41188.000000  41188.000000  41188.000000
mean    40.02406    258.285010    2.567593    962.475454    0.172963
std     10.42125    259.279249    2.770014    186.910907    0.494901
min     17.00000    0.000000    1.000000    0.000000    0.000000
25%     32.00000    102.000000    1.000000    999.000000    0.000000
50%     38.00000    180.000000    2.000000    999.000000    0.000000
75%     47.00000    319.000000    3.000000    999.000000    0.000000
max     98.00000    4918.000000    56.000000    999.000000    7.000000

count    emp.var.rate  cons.price.idx  cons.conf.idx  euribor3m  nr.employed
count  41188.000000    41188.000000    41188.000000  41188.000000  41188.000000
mean     0.081886      93.575664      -40.502600      3.621291    5167.035911
std      1.570960      0.578840       4.628198      1.734447      72.251528
min     -3.400000      92.201000     -50.800000      0.634000    4963.600000
25%     -1.800000      93.075000     -42.700000      1.344000    5099.100000
50%      1.100000      93.749000     -41.800000      4.857000    5191.000000
75%      1.400000      93.994000     -36.400000      4.961000    5228.100000
max      1.400000      94.767000     -26.900000      5.045000    5228.100000

```

Observations : The typical client is around 40 years old.

Calls usually take around 4 to 5 minutes, but some can last for hours.

Clients are usually called 2 or 3 times during a campaign.

As per our problem statement ( Identify key factors that predict if a client will subscribe to a bank's term deposit program.) and by observing data we taking demographic and loan data columns which are the factors for bank deposit programs and not considering other columns

```

[47]: # taking columns which are considered as factors
new_df = df[['age', 'job', 'marital', 'education', 'default', 'housing',
↳ 'loan', 'y']]
new_df

```

```

[47]:
age      job  marital      education  default  housing  loan \
0      56  housemaid  married      basic.4y      no      no      no
1      57  services  married      high.school  unknown      no      no
2      37  services  married      high.school      no      yes      no

```

3	40	admin.	married	basic.6y	no	no	no
4	56	services	married	high.school	no	no	yes
...	...	...	...	...	...	...	...
41183	73	retired	married	professional.course	no	yes	no
41184	46	blue-collar	married	professional.course	no	no	no
41185	56	retired	married	university.degree	no	yes	no
41186	44	technician	married	professional.course	no	no	no
41187	74	retired	married	professional.course	no	yes	no

	y
0	no
1	no
2	no
3	no
4	no

...	...
41183	yes
41184	no
41185	no
41186	yes
41187	no

[41188 rows x 8 columns]

Check the info and then Outliers


```
[8]: new_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 41188 entries, 0 to 41187
Data columns (total 8 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         41188 non-null  int64
1   job         41188 non-null  object
2   marital     41188 non-null  object
3   education   41188 non-null  object
4   default     41188 non-null  object
5   housing     41188 non-null  object
6   loan        41188 non-null  object
7   y           41188 non-null  object
dtypes: int64(1), object(7)
memory usage: 2.5+ MB
```

Cleaning of Data

```
[10]: new_df.age.plot(kind='box')
```

[10]: <Axes: >



We can see many outliers

Re modify our data frame with consideration of all loans

```
[12]: new_df = new_df.rename(columns={
    'y': 'subscription_outcome',
    'marital': 'marital_status',
    'default': 'previously_defaulted_loan',
    'housing': 'has_housing_loan',
    'loan': 'has_personal_loan'})

new_df
```

```
[12]:
```

	age	job	marital_status	education	\
0	56	housemaid	married	basic.4y	
1	57	services	married	high.school	
2	37	services	married	high.school	
3	40	admin.	married	basic.6y	
4	56	services	married	high.school	
...	...	...	...	...	
41183	73	retired	married	professional.course	

41184	46	blue-collar	married	professional.course
41185	56	retired	married	university.degree
41186	44	technician	married	professional.course
41187	74	retired	married	professional.course

	previously_defaulted_loan	has_housing_loan	has_personal_loan	\
0	no	no	no	
1	unknown	no	no	
2	no	yes	no	
3	no	no	no	
4	no	no	yes	
...	...	...	...	
41183	no	yes	no	
41184	no	no	no	
41185	no	yes	no	
41186	no	no	no	
41187	no	yes	no	

	subscription_outcome
0	no
1	no
2	no
3	no
4	no
...	...
41183	yes
41184	no
41185	no
41186	yes
41187	no

[41188 rows x 8 columns]

```
[13]: new_df.age.describe()
```

```
[13]: count    41188.00000
      mean      40.02406
      std       10.42125
      min       17.00000
      25%       32.00000
      50%       38.00000
      75%       47.00000
      max       98.00000
      Name: age, dtype: float64
```

Check value counts



```
[14]: new_df.job.value_counts()
```

```
[14]: job
      admin.          10422
      blue-collar    9254
      technician    6743
      services      3969
      management    2924
      retired       1720
      entrepreneur  1456
      self-employed  1421
      housemaid     1060
      unemployed    1014
      student       875
      unknown       330
      Name: count, dtype: int64
```

```
[15]: new_df.marital_status.value_counts()
```

```
[15]: marital_status
      married    24928
      single     11568
      divorced    4612
      unknown      80
      Name: count, dtype: int64
```

```
[16]: new_df.education.value_counts()
```

```
[16]: education
      university.degree    12168
      high.school          9515
      basic.9y             6045
      professional.course   5243
      basic.4y             4176
      basic.6y             2292
      unknown              1731
      illiterate            18
      Name: count, dtype: int64
```

```
[17]: new_df.previously_defaulted_loan.value_counts()
```

```
[17]: previously_defaulted_loan
      no          32588
      unknown     8597
      yes          3
      Name: count, dtype: int64
```

```
[18]: new_df.has_housing_loan.value_counts()
```

```
[18]: has_housing_loan
yes      21576
no       18622
unknown   990
Name: count, dtype: int64
```

```
[19]: new_df.has_personal_loan.value_counts()
```

```
[19]: has_personal_loan
no      33950
yes      6248
unknown   990
Name: count, dtype: int64
```

for converting categorical to numerical use label encoding

```
[21]: from sklearn.preprocessing import LabelEncoder
label_encoder = LabelEncoder()
```

```
[22]: new_df['job'] = label_encoder.fit_transform(new_df['job'])
new_df.marital_status = label_encoder.fit_transform(new_df.marital_status)
new_df.education = label_encoder.fit_transform(new_df.education)
new_df.previously_defaulted_loan = label_encoder.fit_transform(new_df.
    ↳previously_defaulted_loan)
new_df.has_housing_loan = label_encoder.fit_transform(new_df.has_housing_loan)
new_df.has_personal_loan = label_encoder.fit_transform(new_df.has_personal_loan)
new_df.subscription_outcome = label_encoder.fit_transform(new_df.
    ↳subscription_outcome)
```

```
[23]: new_df
```

```
[23]:
```

	age	job	marital_status	education	previously_defaulted_loan	\
0	56	3	1	0		0
1	57	7	1	3		1
2	37	7	1	3		0
3	40	0	1	1		0
4	56	7	1	3		0
...	...	...	...	...	...	
41183	73	5	1	5		0
41184	46	1	1	5		0
41185	56	5	1	6		0
41186	44	9	1	5		0
41187	74	5	1	5		0

```
has_housing_loan has_personal_loan subscription_outcome
```

```

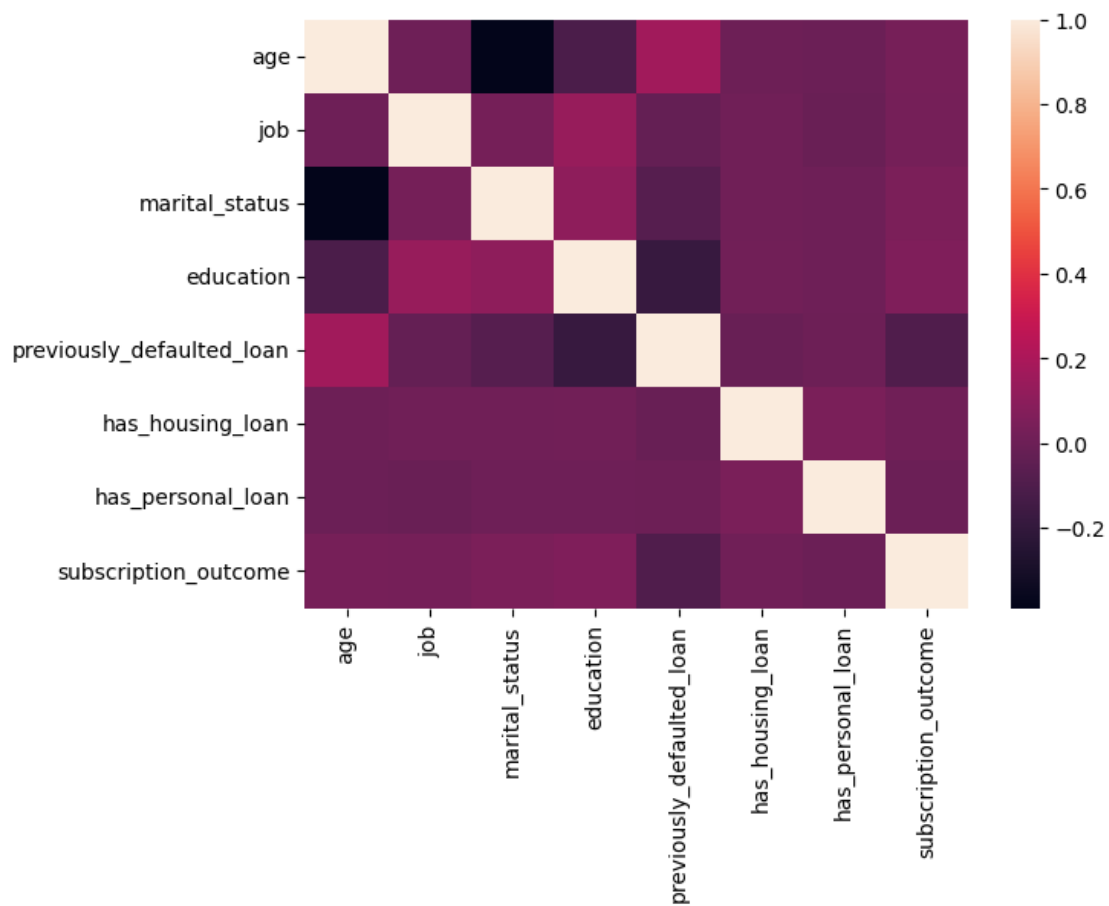
0          0          0          0
1          0          0          0
2          2          0          0
3          0          0          0
4          0          2          0
...
41183      ...          ...          1
41184      0          0          0
41185      2          0          0
41186      0          0          1
41187      2          0          0

```

[41188 rows x 8 columns]

Plot a heat map

```
[29]: plt.figure(figsize=(7,5))
      _ = sns.heatmap(new_df.corr())
```



There are some relationships between two features eg age and previously\_defaulted\_loan, education and job etc, but the relationships are not very strong (you can clearly observe from heat map)

**For creating Model, Select features and Split the dataset into training and test sets**

```
[36]: # Assume 'new_df' is your DataFrame and the last column 'y' is the target
      ↪variable

      # Select features (all columns except the last one) and the target (the last
      ↪column)
      X = new_df.iloc[:, :-1]
      y = new_df.iloc[:, -1]

      # Split the dataset into training and test sets
      X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
      ↪random_state=42)
```

```
[37]: model = KNeighborsClassifier(n_neighbors=4, n_jobs=-1)
      model.fit(X_train, y_train)
      score = model.score(X_test, y_test)
      score
```

```
[37]: 0.8816460305899491
```

we got a pretty good model score, This means that even though the details we have don't seem to relate to each other very much, we can still combine them to make a model that can guess if a client might subscribe for loan

Check accuracy (i.e Optimal value of K)with different number of neighbours

```
[38]: accuracy_score = []

      for k in range(1, 11):
          model = KNeighborsClassifier(n_neighbors=k, n_jobs=-1)
          model.fit(X_train, y_train)
          score = model.score(X_test, y_test)
          accuracy_score.append(score)

      print(f'K value {k}: {score}')
```

```
K value 1: 0.8159747511531925
K value 2: 0.8771546491866958
K value 3: 0.8622238407380433
K value 4: 0.8816460305899491
K value 5: 0.872784656470017
K value 6: 0.8815246419033747
K value 7: 0.8792182568584608
K value 8: 0.8839524156348628
```


K value 9: 0.8826171400825443  
K value 10: 0.8844379703811605

```
[39]: max(accuracy_score)
```

```
[39]: 0.8844379703811605
```

Max accuracy is observed at k=4 and see the plot for all varying of K

```
[42]: plt.scatter(range(1, 11), accuracy_score);  
plt.plot(range(1, 11), accuracy_score);  
plt.xticks(range(1, 11));
```



**Making Your KNN Model Better** To make our KNN model better, we can:

1. Adjust settings like how many neighbors the model should consider. Changing these can make the model predict better, but just adding more neighbors doesn't always help.
2. Cut down on less important data but keep the crucial stuff. Our data is already pretty trimmed, so cutting more might lose important bits.
3. Make sure all the data is on a similar scale so that no single thing overpowers the rest. This can help since things like age and job are on different scales.

We'll try scaling our data to see if our model gets better.

```
[ ]: Use StandardScaler and compare results with KNN accuracies
```

```
[43]: norm_trained_model_scores = []
scaled_dataTrained_model = []
for k in range(1, 11):
    data_scaled_model = Pipeline([('Scaler', StandardScaler()), ('KNN',
↪KNeighborsClassifier(n_neighbors=k, n_jobs=-1))])
    data_scaled_model.fit(X_train, y_train)
    data_scaled_model_score = data_scaled_model.score(X_test, y_test)
    scaled_dataTrained_model += [data_scaled_model_score]

    norm_trained_model = KNeighborsClassifier(n_neighbors=k, n_jobs=-1)
    norm_trained_model.fit(X_train, y_train)
    norm_trained_model_score = norm_trained_model.score(X_test, y_test)
    norm_trained_model_scores += [norm_trained_model_score]

    print(f'K value {k}: scaled data model accuracy score:
↪{data_scaled_model_score} | Normal data mode accuracu score:
↪{norm_trained_model_score}')
```

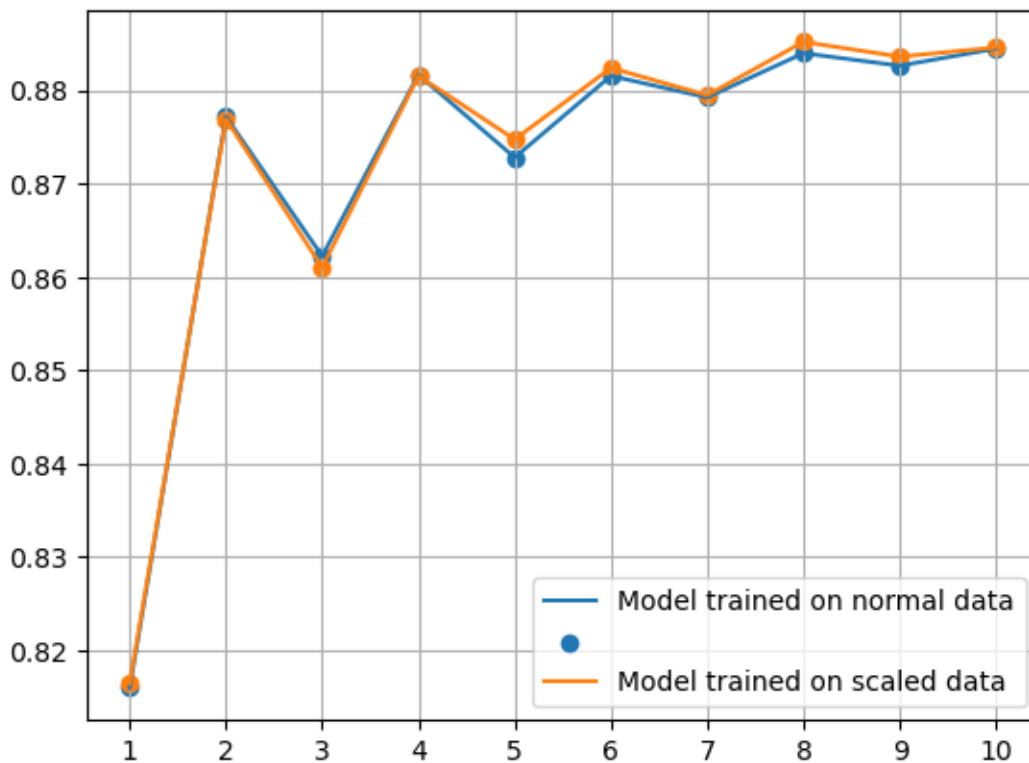
```
K value 1: scaled data model accuracy score: 0.8164603058994901 | Normal data
mode accuracu score: 0.8159747511531925
K value 2: scaled data model accuracy score: 0.876911871813547 | Normal data
mode accuracu score: 0.8771546491866958
K value 3: scaled data model accuracy score: 0.8610099538722991 | Normal data
mode accuracu score: 0.8622238407380433
K value 4: scaled data model accuracy score: 0.8815246419033747 | Normal data
mode accuracu score: 0.8816460305899491
K value 5: scaled data model accuracy score: 0.8747268754552076 | Normal data
mode accuracu score: 0.872784656470017
K value 6: scaled data model accuracy score: 0.8823743627093955 | Normal data
mode accuracu score: 0.8815246419033747
K value 7: scaled data model accuracy score: 0.8794610342316096 | Normal data
mode accuracu score: 0.8792182568584608
K value 8: scaled data model accuracy score: 0.8851663025006069 | Normal data
mode accuracu score: 0.8839524156348628
K value 9: scaled data model accuracy score: 0.8835882495751396 | Normal data
mode accuracu score: 0.8826171400825443
K value 10: scaled data model accuracy score: 0.8845593590677349 | Normal data
mode accuracu score: 0.8844379703811605
```

```
[45]: # display the results
plt.plot(range(1, 11), norm_trained_model_scores)
plt.scatter(range(1, 11), norm_trained_model_scores)

plt.plot(range(1, 11), scaled_dataTrained_model)
plt.scatter(range(1, 11), scaled_dataTrained_model)

plt.grid()
plt.xticks(range(1, 11))
plt.legend(['Model trained on normal data', None, 'Model trained on scaled_
↪data'])
```

[45]: <matplotlib.legend.Legend at 0x2203d361dd0>



Looking at how well the two models did, we can tell they're both doing well. they're at least 80% accurate. The best models are more than 88.44% accurate.

**conclusion** Our model shows key factors that predict if a client will subscribe to a bank's term deposit program. just by using basic information about them and their loan history. It's impressive that our model reached an accuracy of 88.44% with this limited information.

To make the model even better, we need to scale the data. This means adjusting the features so they're all measured the same way. Doing this helps prevent any single feature from having too

much influence and can make our predictions more accurate.

[ ]:

[ ]: